

Replicating *Political Deepfakes* (Barari, Lucas, and Munger (2025))

Zoe Meers

September 21, 2026

This research project replicates Barari, Lucas, and Munger (2025) on partisan deepfake credibility and detection using a pre-registered randomized control trials (RCT) with $n = 5,724$ participants conducted on survey platform Lucid. In their work, (Barari, Lucas, and Munger 2025) find partisan effects on deepfake detection accuracy when their in-party political leader is involved in a political scandal.¹ Barari, Lucas, and Munger (2025) develop two sequential experimental designs to test deepfake video exposure and detection accuracy:

1. In the first exposure experiment, the authors implement a 2x6 factorial design where deepfakes were mixed in with other media messaging (text, audio, or a political skit) involving 2020 Democratic primary candidate Elizabeth Warren. The design and layout of the experiment mimics a Facebook newsfeed, where the incorrect content is surrounded by two (real) pieces of media before and after the deepfake video.
2. In the second experiment, respondents are asked to discriminate between authentic and deepfake videos. A list of eight videos, split into three groupings (no deepfakes, low rate of deepfakes (2/8 videos are deepfakes), and high rate of deepfakes (6/8 are deepfakes)), is randomly assigned to participants.

I plan to replicate the second experiment focusing on deepfake detection.² I would like to extend this work by considering the false-positive rate in more depth. I am particularly interested in testing the following hypothesis:

- H1: Political partisans have different false positive rates in the no-deepfake condition, showing diverging partisan skepticism about the authenticity of deepfake videos.³

I am also interested in providing further analysis on heterogeneous treatment effects. I plan to use a flexible machine learning method Bayesian Additive Regression Trees (as demonstrated in by (Green and Kern 2012)) to look beyond fixed subsets in CATEs. This will be a secondary analysis.

¹That is, partisans are more likely to think a real video is a deepfake when the scandal refers to someone in their own party. This is also defined as the false positive rate.

²My replication files will be available here: <https://github.com/zmeers/political-deepfake-fx>.

³In other words, strong Republicans or Democrats do not start off on a level-playing field regarding deepfake learning.

References

- Barari, Soubhik, Christopher Lucas, and Kevin Munger. 2025. “Political Deepfakes Are as Credible as Other Fake Media and (Sometimes) Real Media.” *The Journal of Politics* 87 (2): 510–26.
- Green, Donald P, and Holger L Kern. 2012. “Modeling Heterogeneous Treatment Effects in Survey Experiments with Bayesian Additive Regression Trees.” *Public Opinion Quarterly* 76 (3): 491–511.